

# GSoC 2022: Remove undef: Move Uninitialized Memory to Poison Final Report

John McIver (mentee)  
mciverj@unm.edu

Nuno Lopes (mentor)  
nuno.lopes@tecnico.ulisboa.pt

September 8, 2022

## Abstract

Current memory semantics of the LLVM-IR use `undef` when dereferencing an uninitialized memory location. This results in `phi` nodes comprised of `undef` operands (instead of `poison`), which prevents further optimizations, such as the folding of `phi` nodes. The solution is to convert uninitialized memory location loads to `poison`. This Google Summer of Code (GSoC) project implements methodologies for bit-field loading and storing and replacing `phi-undef` with `phi-poison` or `phi-freeze-poison` for conditional stores in SROA, `mem2reg`, and GVN. This work increases the use of `poison` semantics, contributing to the eventual deprecation of `undef`.

## 1 Introduction

During the period of performance the following topics were developed for community review and patch submission:

- Refactor clang IR generation of C and C++ bit-fields to utilize `poison` memory semantics
- Refactor memory based optimizations to use `phi` with `poison` rather than `phi` with `undef`

The refactoring of memory based optimizations was expanded from the original proposal beyond scalar replacement of aggregates (SROA) to include `mem2reg` and global value numbering (GVN), NewGVN, and Glob-

alOpt. Additionally, refactoring of memory state built-ins used by several optimizations has been refactored.

Prior to the beginning of GSoC I had limited experience with the LLVM framework and patch submission process [4, 3]. During the implementation of this project I gained experience with:

- LLVM optimization framework
- LLVM IR
  - Semantics of `undef`, `freeze poison`, etc.
  - Attribute implementation and adaption
  - Memory access semantics
  - Data structure querying and manipulation
- LLVM regression and test-suite frameworks
- Performance testing and root cause analysis techniques
- Automatic test case reduction

Gaining experience with these topics has greatly improved by understanding of the project's architecture and my ability to identify and develop solutions within the LLVM framework.

## 2 Bit-field Migration to Poison

A request for comment (RFC) was developed and published on the LLVM Discourse site [6], which outlines three different solutions: emit freeze on all bit-field loads, zero initialize and forego freeze on stack allocated bit-fields, and freeze based on access analysis. While the proposal was under going community review a patch was developed using the emission of freeze on all bit-field loads [2].

Bootstrap builds of the test-suite showed that the patch [2] resulted in as much a 15% increase in runtime for `SIBsim4`. Analysis revealed that the addition of `freeze` instructions in bit-fields resulted in extra `mov` instructions. Refactoring code generation to remove the extraneous `mov` instructions was attempted, however due to time constraints an issues was filed [7].

Two issues were discovered during the RFC process. First, an aliasing scheme would need to be introduced in order to allow bit-field freezes based on access analysis. Second, inconsistent use of option

`-ffine-grained-bitfield-accesses` can cause incorrect behavior when linking with inter-procedural optimization (IPO). To reduce the risk of IPO fine-grained bit-field inlining interactions, two patches were developed to add an IR function attribute `find.fine-grained-bitfields` [5, 1]. Patch [5] was rejected as it introduces language specific knowledge at the IR level.

### 3 Load with Automatic Freeze

Current `undef` usage appears around dead operations where loads of uninitialized memory occur. A preliminary RFC was posted to LLVM Discourse outlining three potential solutions [8]:

	Can duplicate?	Can hoist?
1. Load with automatic freeze	No	Yes
2. Load without automatic freeze	Yes	Yes
3. Load with <code>noundef</code>	Yes	No

Currently Clang generally emits option 3, which can also be downgraded to option 2 if hoisting is required. Option 1 was selected for initial implementation.

The first optimizations to be augmented with the automatic freeze poison semantics was SROA and `mem2reg`. The optimization `mem2reg` is used by SROA to promote `alloca` instructions. Unlike the insertion of `undef` constants the insertion of freeze poison instructions requires tracking for potential removal.

GVN adaptation has also been implemented, but with modification to the memory built-ins utility function `getInitialValueOfAllocation`, which now returns poison on uninitialized memory. This API is currently undergoing refactoring to provide two methods: `is initialized to zero` and `is uninitialized`. Similar to `mem2reg`, GVN partial redundancy elimination (PRE) requires the tracking of freeze poison insertions for later removal due to critical edge incompatibilities.

Support for the emission of poison in the presence of uninitialized load instructions in the presence of attribute `!noundef` has only been implemented in `mem2reg`. [SROA should be completed this week.](#)

### 4 Future Work

Several patches are currently outstanding due to related technical issues. The mentee has agreed to continue work on the project past the period of

performance with support from the mentor. The following sections outline the outstanding issues and possible resolutions.

#### 4.1 Load with Automatic Freeze

A formal RFC for Load with Automatic Freeze (see Section 3) needs to be submitted on Discourse. Modification and patch submission will need to occur for mem2reg, SROA, GVN, and associated LTO/IPO optimizations. Analysis will be performed to verify that load duplication is not occurring and if its appropriate mitigation techniques will be developed. If during RFC review major technical issues are discovered then alternative semantics will be developed and implemented as per community input.

### 5 Conclusion

This project has provided a great opportunity to learn about many of the components of LLVM. This has greatly increased my understanding of IR semantics, compiler frameworks, and general compiler development methodologies. GSoC has provided a great start and I look forward to seeing this project to completion.

### References

- [1] John McIver. 2022. [codegen] add codegen of ir function attribute fine\_grained\_bitfields. (July 2022). <https://reviews.llvm.org/D129542>.
- [2] John McIver. 2022. [codegen] make uninitialized lvalue bit-field stores poison compatible. (June 2022). <https://reviews.llvm.org/D128501>.
- [3] John McIver. 2022. [instcombine] optimize and of icmps with power-of-2 and contiguous masks. (May 2022). <https://reviews.llvm.org/D125717>.
- [4] John McIver. 2022. [llvm-objcopy] llvm-strip option -only-keep-debug should suppress default -strip-all. (April 2022). <https://reviews.llvm.org/D123798>.
- [5] John McIver. 2022. [llvm][ipo] add ir function attribute fine\_grained\_bitfields. (July 2022). <https://reviews.llvm.org/D129541>.

- [6] John McIver. 2022. [rfc] making bit-field codegen poison compatible. (June 2022). <https://discourse.llvm.org/t/rfc-making-bit-field-codegen-poison-compatible/63250>.
- [7] John McIver. 2022. Llvm codegen producing extra move instructions or crashing with freeze #56646. (July 2022). <https://github.com/llvm/llvm-project/issues/56646>.
- [8] John McIver. 2022. Remove undef: move uninitialized memory to poison. (July 2022). <https://discourse.llvm.org/t/remove-undef-move-uninitialized-memory-to-poison/61123>.